

# Neural Network Function Approximation

# Map of RL Algorithms

Model Based  
←

Model Free  
→

Learn $Q$	Learn $\pi_\theta$
SARSA	Policy Gradient

On Policy

Q-learning

Off Policy

MLMBTRL  
(Learn  $T, R$ )

Tabular

# This Time

Challenges in Reinforcement Learning:


- Exploration vs Exploitation ← Bandit
- Credit Assignment ←
- Generalization ← Today

# Function Approximation

# Function Approximation

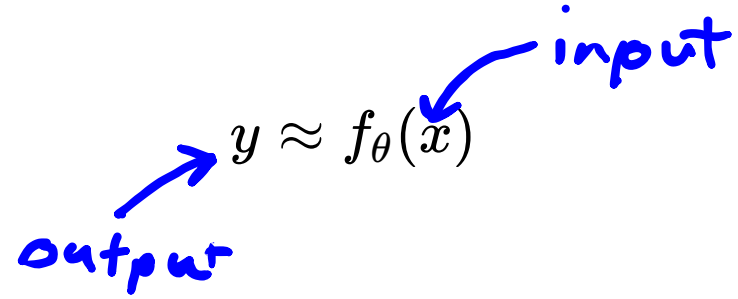
$$y \approx \underset{\tau}{f_{\theta}}(x)$$

# Function Approximation

$$y \approx f_{\theta}(x)$$


input

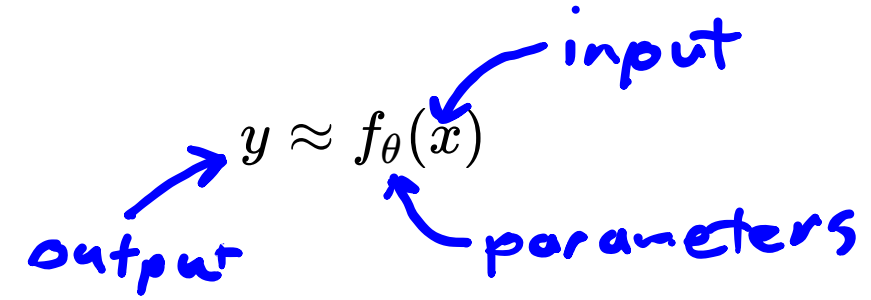
# Function Approximation



A diagram illustrating the function approximation equation  $y \approx f_{\theta}(x)$ . The equation is written in black. Two blue arrows point to the variables: one from the word "output" to  $y$ , and another from the word "input" to  $x$ . The words "output" and "input" are written in a blue, handwritten style.

$$y \approx f_{\theta}(x)$$

# Function Approximation



A diagram illustrating the function approximation equation  $y \approx f_{\theta}(x)$ . The equation is centered, with three handwritten blue arrows pointing to its components: one from the word "output" to  $y$ , one from the word "input" to  $x$ , and one from the word "parameters" to  $\theta$ .

$$y \approx f_{\theta}(x)$$

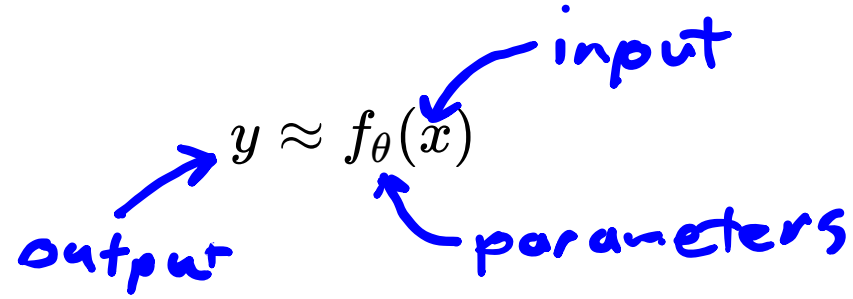
output

input

parameters



# Function Approximation

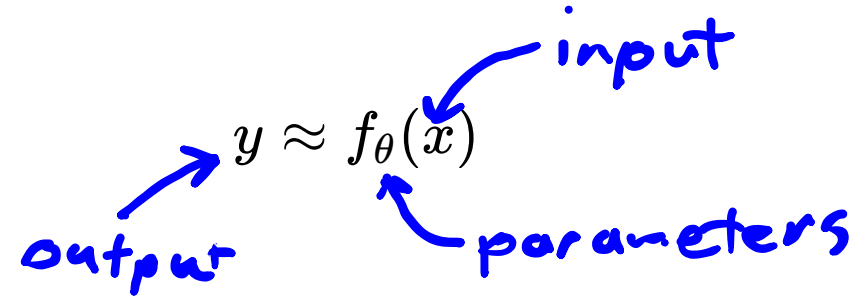


A diagram illustrating the function approximation equation  $y \approx f_{\theta}(x)$ . The equation is centered, with three handwritten blue annotations: an arrow pointing from the word "output" to  $y$ , an arrow pointing from the word "input" to  $x$ , and an arrow pointing from the word "parameters" to  $\theta$ .

Previously, Linear:

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

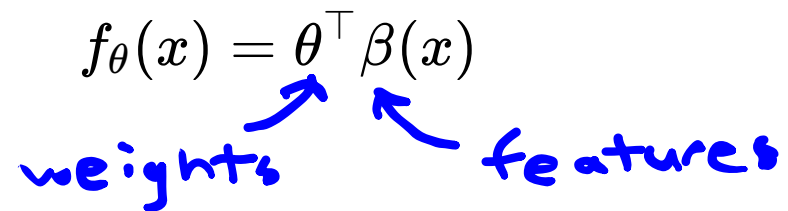
# Function Approximation



A diagram showing the general function approximation equation  $y \approx f_{\theta}(x)$ . Three blue arrows point to the components: one from the word "output" to  $y$ , one from the word "input" to  $x$ , and one from the word "parameters" to  $\theta$ .

$$y \approx f_{\theta}(x)$$

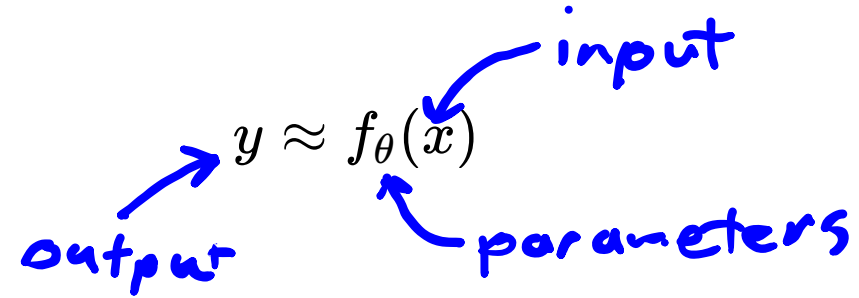
Previously, Linear:



A diagram showing the linear function approximation equation  $f_{\theta}(x) = \theta^{\top} \beta(x)$ . Two blue arrows point to the components: one from the word "weights" to  $\theta$ , and one from the word "features" to  $\beta$ .

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

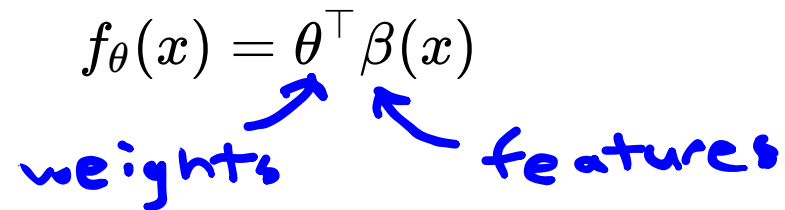
# Function Approximation



A diagram showing the general function approximation equation  $y \approx f_{\theta}(x)$ . Three blue arrows point to the components: one from the word "output" to  $y$ , one from the word "input" to  $x$ , and one from the word "parameters" to  $\theta$ .

$$y \approx f_{\theta}(x)$$

Previously, Linear:



A diagram showing the linear function approximation equation  $f_{\theta}(x) = \theta^{\top} \beta(x)$ . Two blue arrows point to the components: one from the word "weights" to  $\theta$ , and one from the word "features" to  $\beta$ .

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

e.g.  $\beta_i(x) = \sin(i \pi x)$

# Neural Network

# Neural Network



$$h(x) = \sigma(Wx + b)$$

# Neural Network

$$f_{\theta}(x)$$

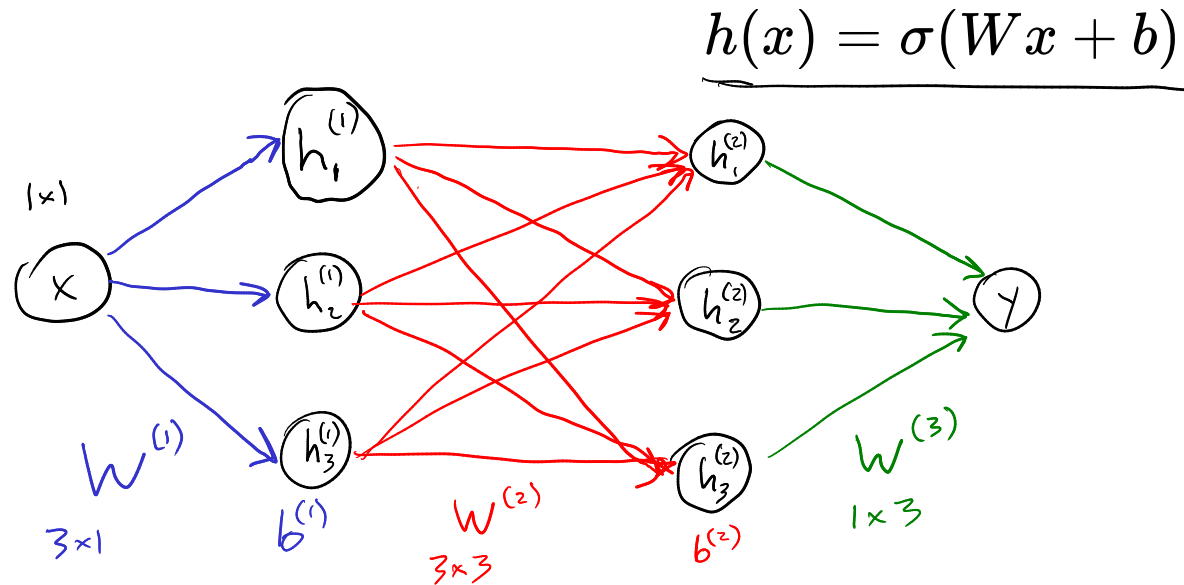
$$h_3(h_2(h_1(x)))$$

$$h(x) = \sigma(Wx + b)$$



# Neural Network

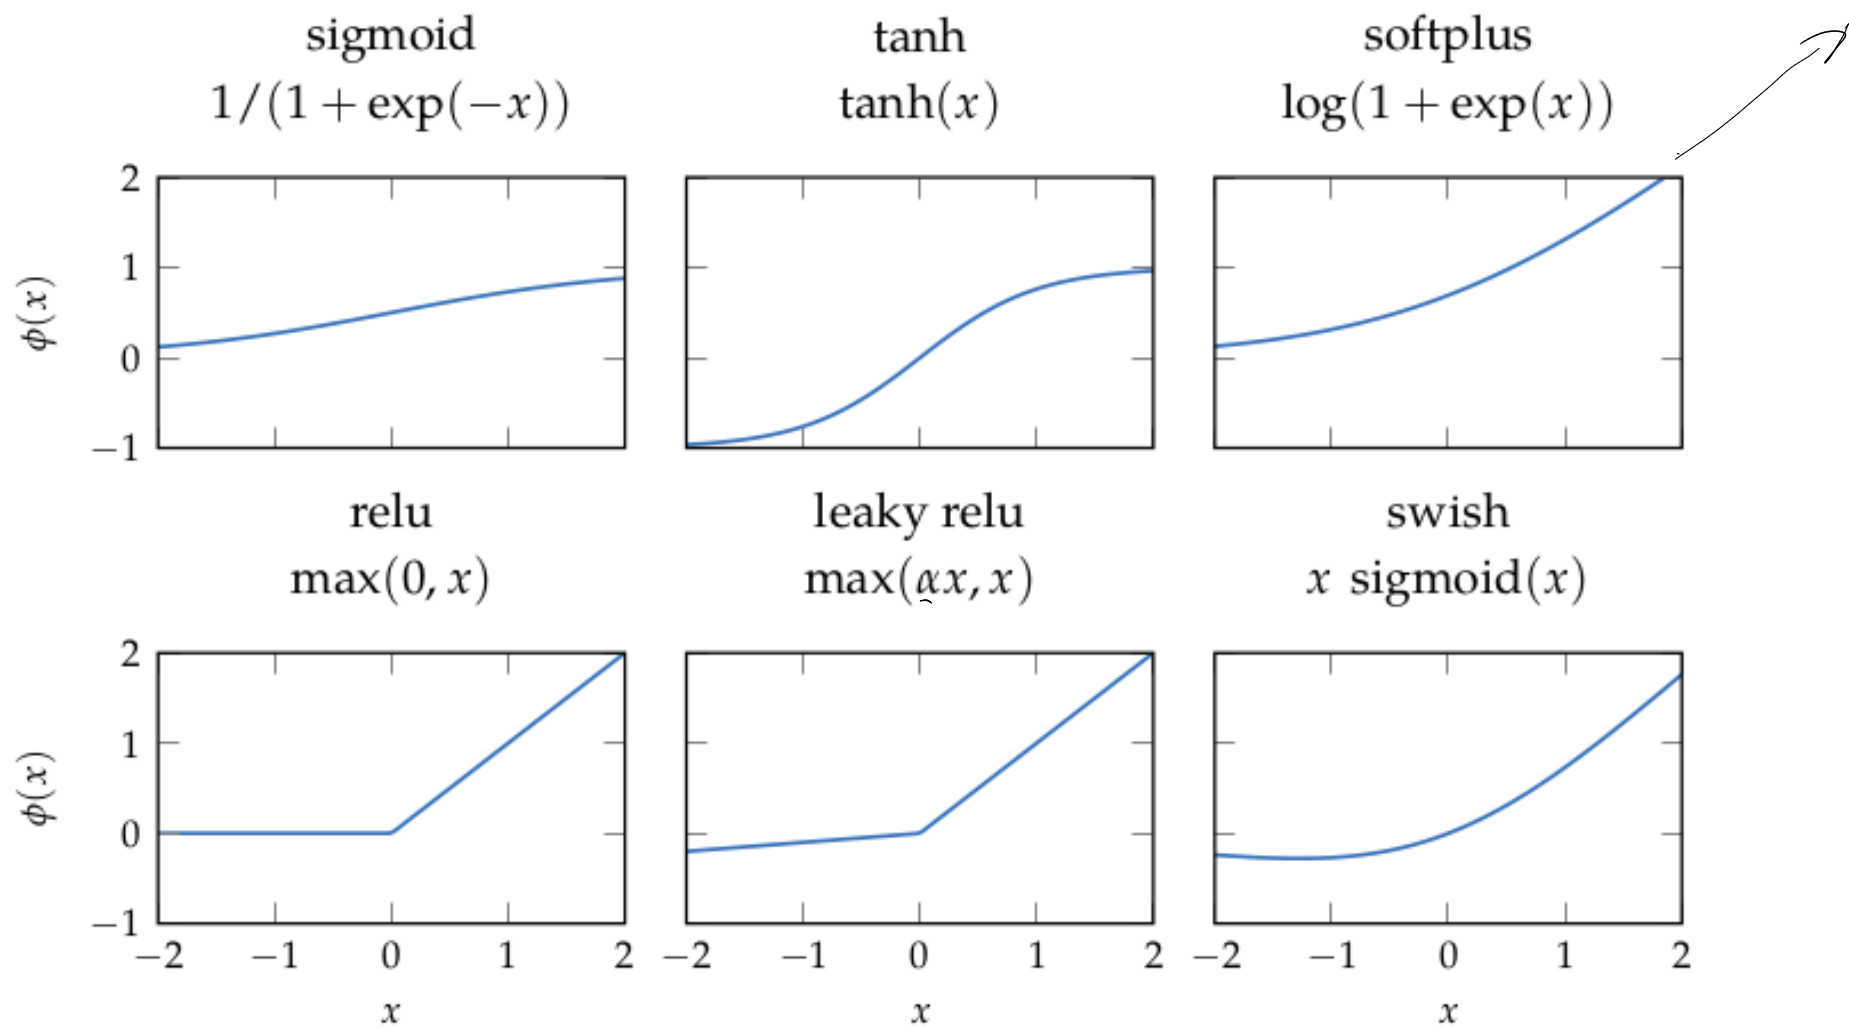
# Neural Network



$$\underline{f_{\theta}(x)} = \underline{h^{(2)}(h^{(1)}(x))} = \underline{W^{(3)}} \underline{\sigma^{(2)}} \left( \underline{W^{(2)}} \underline{\sigma^{(1)}} \left( \underline{W^{(1)}} x + \underline{b^{(1)}} \right) + \underline{b^{(2)}} \right)$$




# Nonlinearities




# Training


# Training



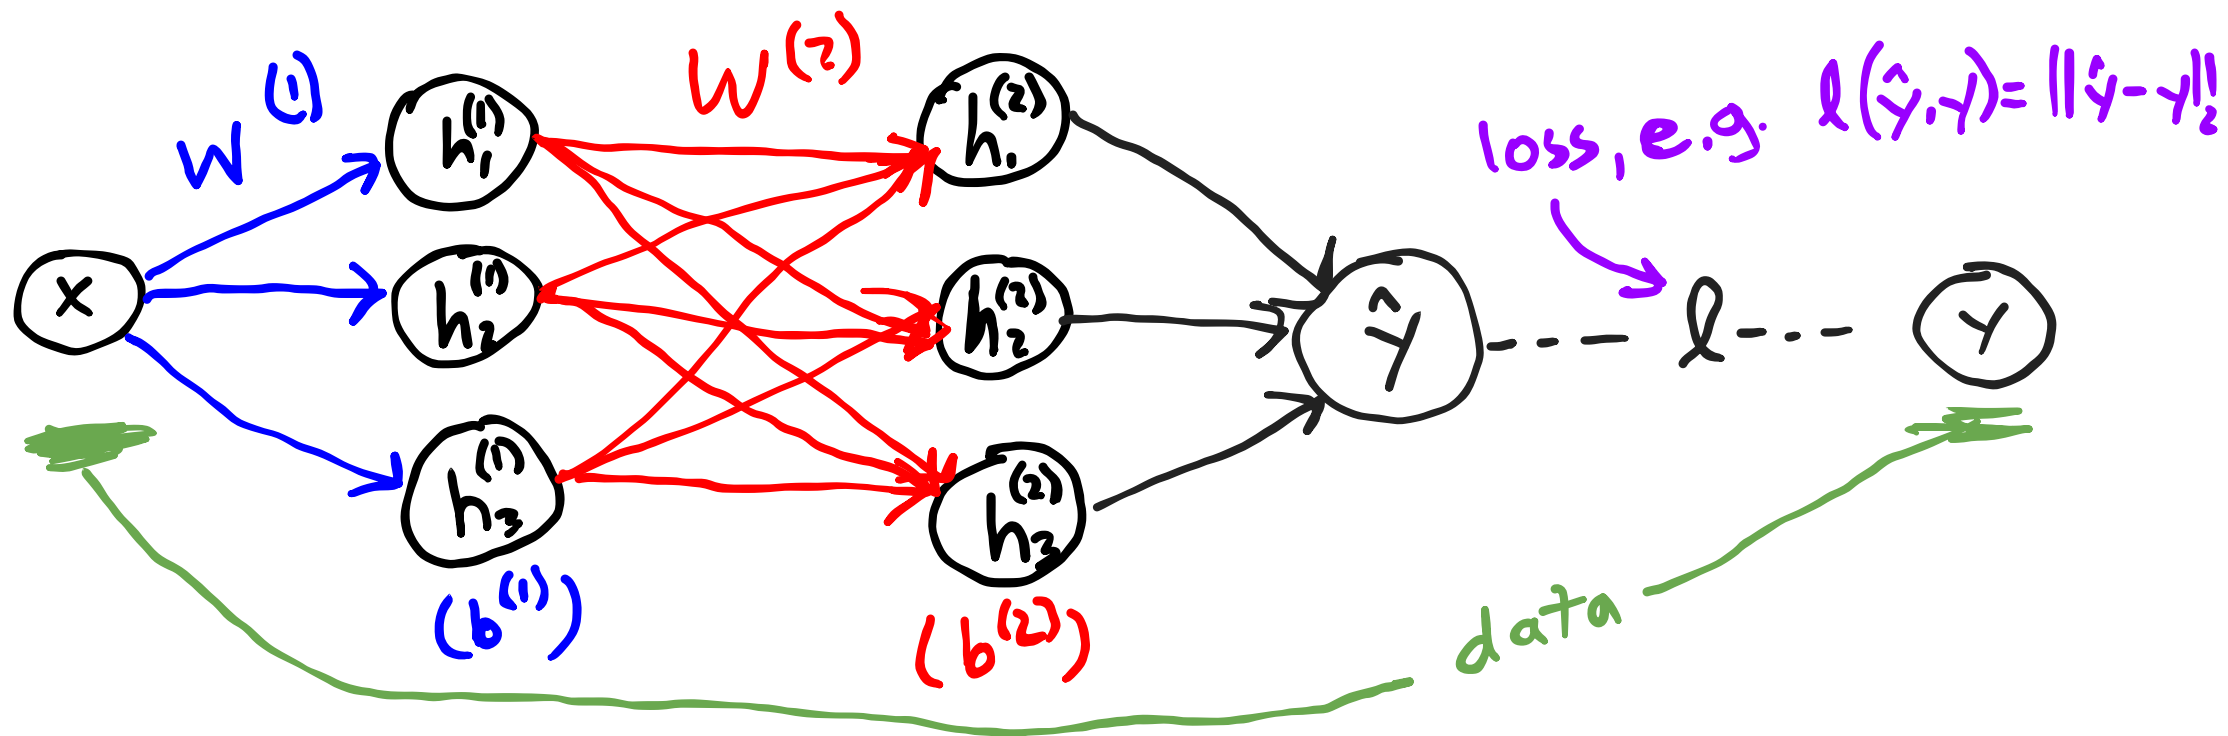
# Training




# Training



# Training

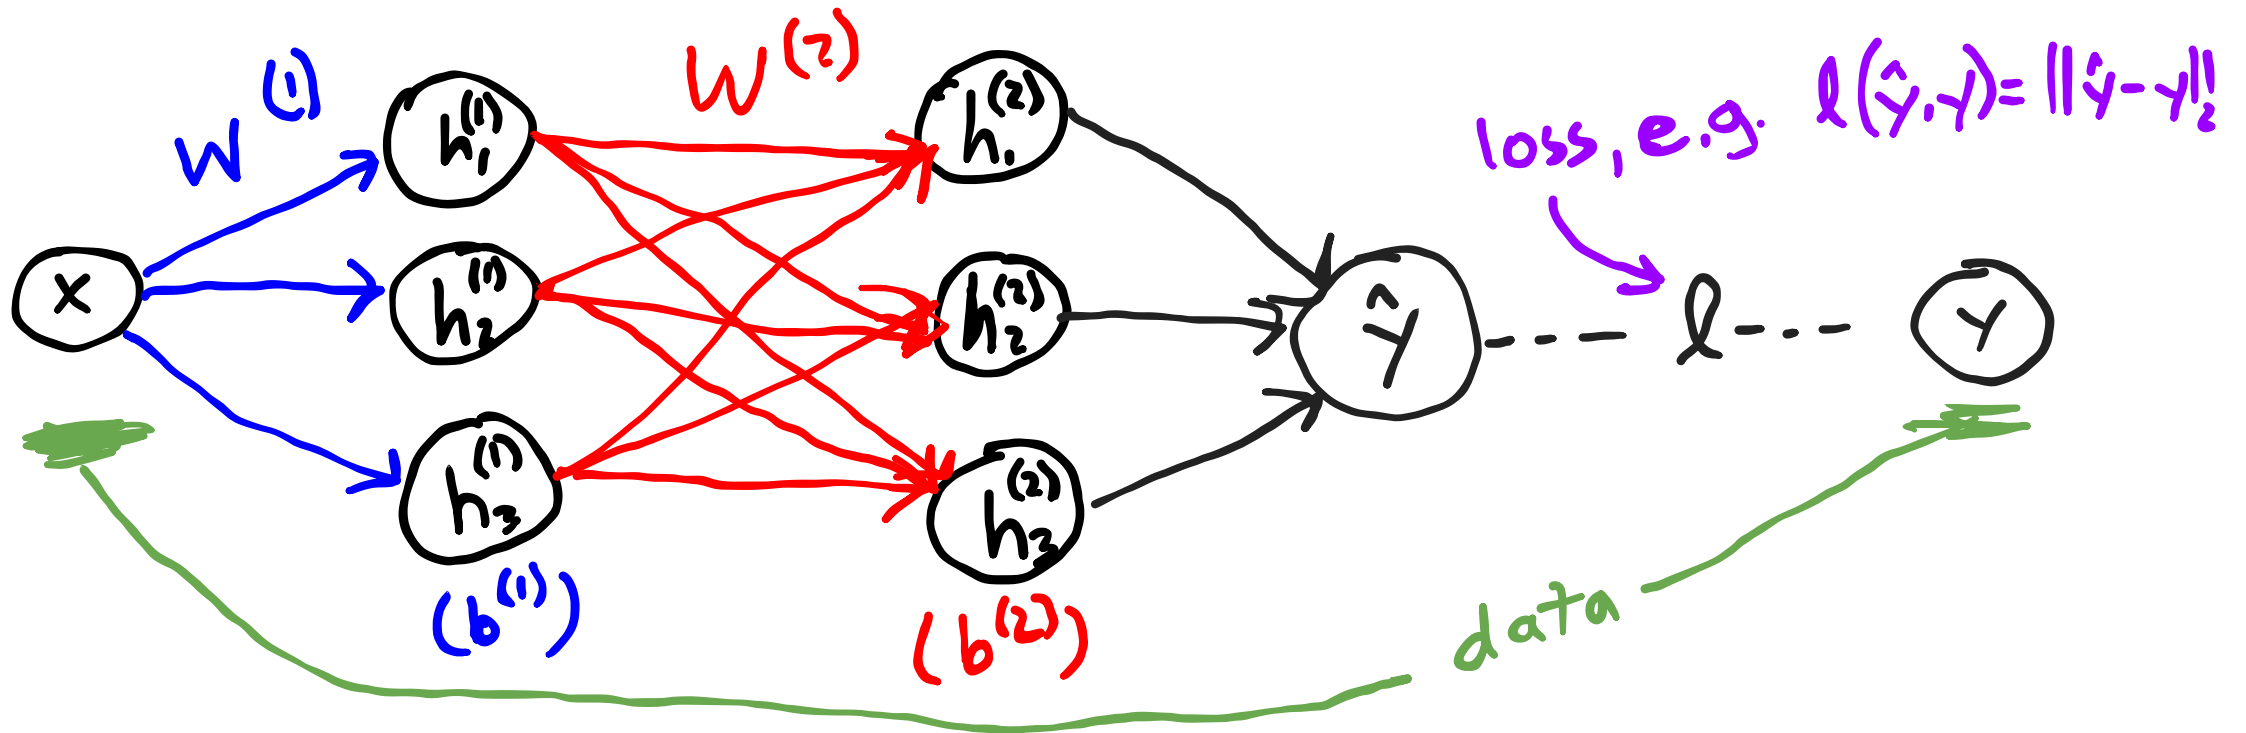


# Training



$$\theta^* = \arg \min_{\theta} \sum_{(x,y) \in \mathcal{D}} l(f_{\theta}(x), y)$$

# Training



$$\theta^* = \arg \min_{\theta} \sum_{(x,y) \in \mathcal{D}} l(f_{\theta}(x), y)$$

Stochastic Gradient Descent:  $\theta \leftarrow \theta - \alpha \nabla_{\theta} l(f_{\theta}(x), y)$



# Chain Rule

$$h_0 = h(x_0)$$

$$f \circ g \circ h = f(g(h(\cdot)))$$

$$\begin{aligned} \left. \frac{\partial f \circ g \circ h}{\partial x} \right|_{x_0} &= \left. \frac{\partial f(g(h(x)))}{\partial x} \right|_{x_0} = \left. \frac{\partial f(g(h))}{\partial h} \right|_{h_0} \left. \frac{\partial h}{\partial x} \right|_{x_0} \\ &= \left. \frac{\partial f(g)}{\partial g} \right|_{g_0} \left. \frac{\partial g(h)}{\partial h} \right|_{h_0} \left. \frac{\partial h}{\partial x} \right|_{x_0} \end{aligned}$$

$$\nabla_{\theta} l(f_{\theta}(x), y)$$

$$\begin{bmatrix} \frac{\partial l}{\partial \theta_1} \\ \vdots \\ \frac{\partial l}{\partial \theta_n} \end{bmatrix}$$

$$\theta = (w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)})$$

$$l(\hat{y}, y) = \frac{1}{n} \sum_i (\hat{y}_i - y_i)^2$$

$$\hat{y} = w^{(2)} \sigma(w^{(1)} x + b^{(1)}) + b^{(2)}$$


$$\left. \frac{\partial l}{\partial w^{(2)}} \right|_{x_0} = \left. \frac{\partial l}{\partial \hat{y}} \right|_{\hat{y}_0} \left. \frac{\partial \hat{y}}{\partial w^{(2)}} \right|_{x_0} = \left. \frac{\partial l}{\partial \hat{y}} \right|_{\hat{y}_0} \left( \sigma(w^{(1)} x_0 + b^{(1)}) \right)$$

# Backprop

# Backprop

$$\hat{y} = \theta_2 x + \theta_1$$

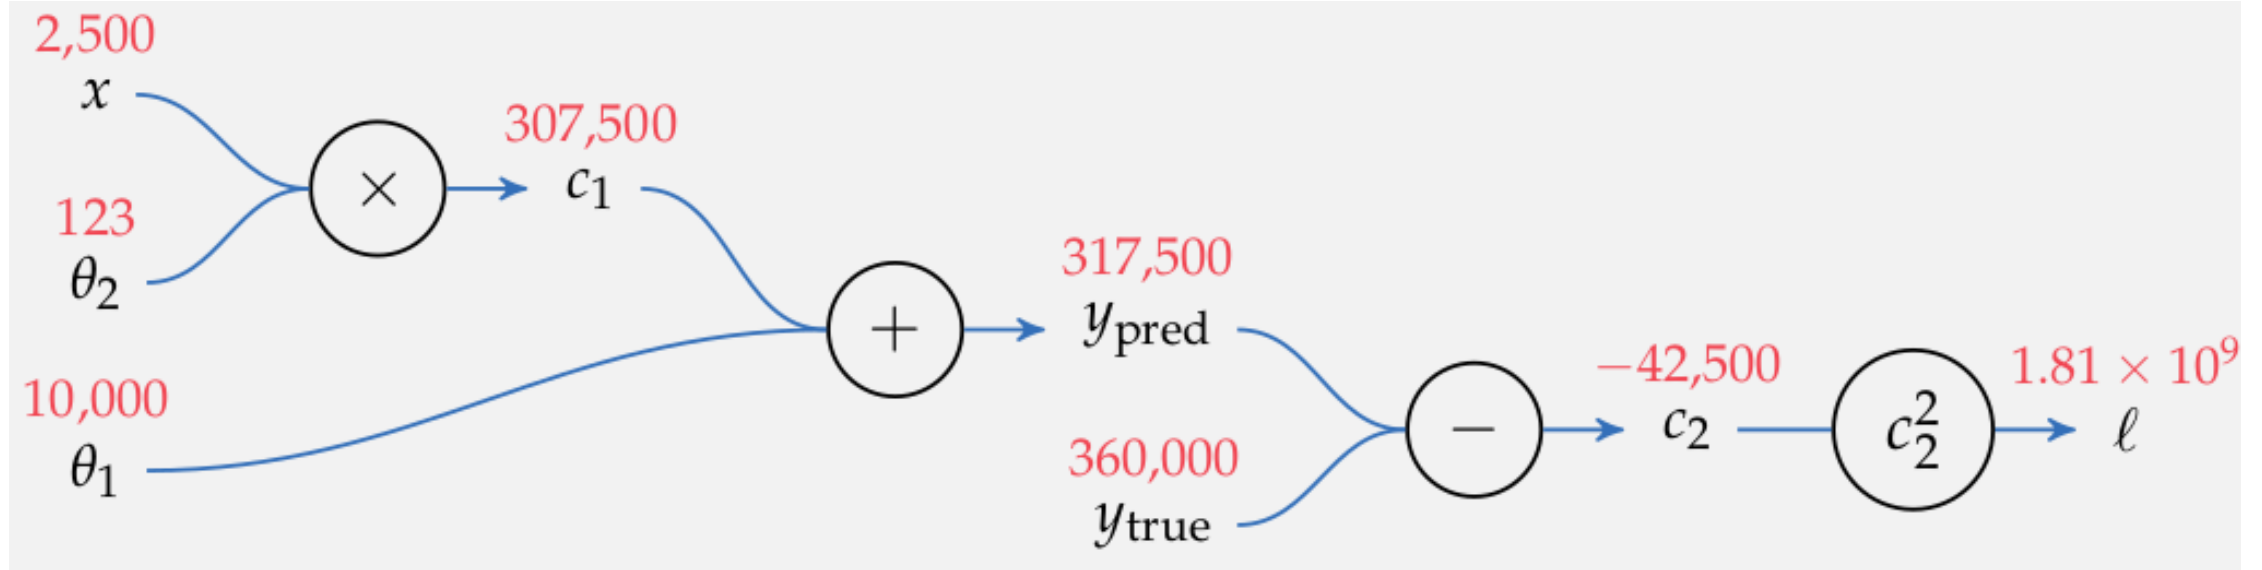
$$l(\hat{y}, y) = (\hat{y} - y)^2$$



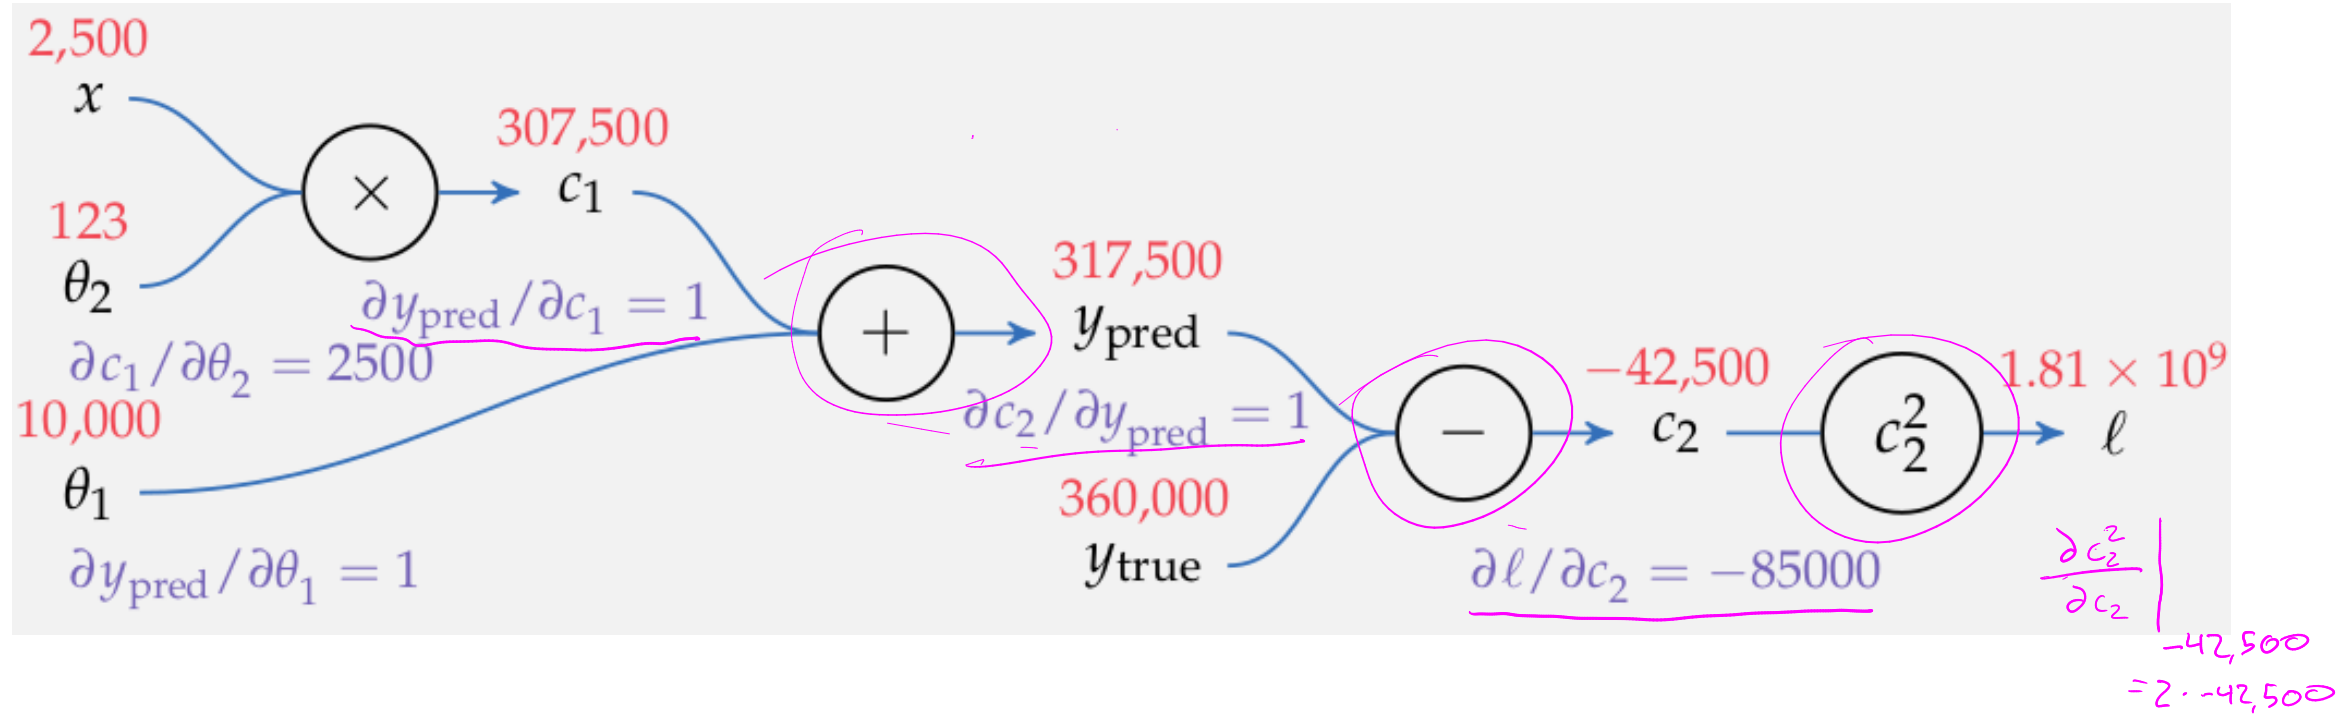
forward: calculate values

backward: use chain rule to calculate derivatives

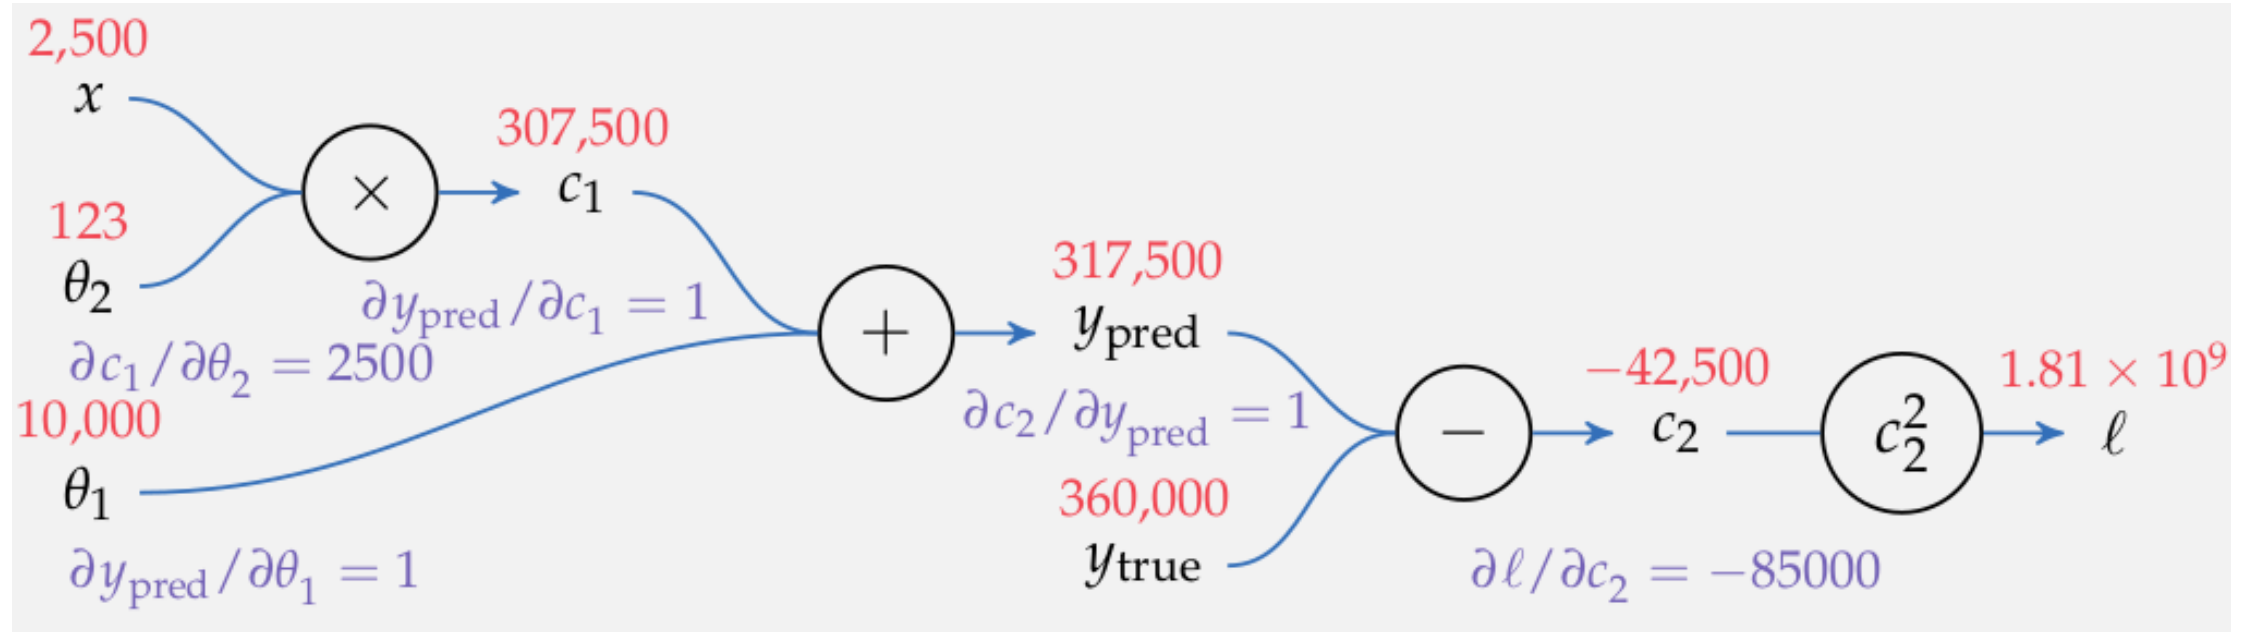
# Backprop



# Backprop



# Backprop



backwards from  $\ell$  to  $\theta_i$

$$\nabla_{\theta} \ell \begin{cases} \frac{\partial \ell}{\partial \theta_1} = \frac{\partial \ell}{\partial c_2} \frac{\partial c_2}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial \theta_1} = -85,000 \cdot 1 \cdot 1 = -85,000 \\ \frac{\partial \ell}{\partial \theta_2} = \frac{\partial \ell}{\partial c_2} \frac{\partial c_2}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial c_1} \frac{\partial c_1}{\partial \theta_2} = -85,000 \cdot 1 \cdot 1 \cdot 2,500 = -2.125 \times 10^8 \end{cases}$$

a “fast and furious” approach to training neural networks does not work and only leads to suffering. Now, suffering is a perfectly natural part of getting a neural network to work well, but it can be mitigated by being thorough, defensive, paranoid, and obsessed with visualizations of basically every possible thing. The qualities that in my experience correlate most strongly to success in deep learning are patience and attention to detail.

Keep calm and  
lower your learning  
rate

- Andrej Karpathy

# Adaptive Step Size: RMSProp

SGD

$$\theta \leftarrow \theta - \alpha \underbrace{\nabla_{\theta} l(f_{\theta}(x), y)}_{g^{(k)}}$$

RMS prop

estimate  
of  $g \odot g$

$$\hat{s}^{(k+1)} \leftarrow \gamma \hat{s}^{(k)} + (1-\gamma) \underbrace{(g^{(k)} \odot g^{(k)})}$$


$$\theta_i^{(k+1)} = \theta_i^{(k)} - \frac{\alpha}{\epsilon + \sqrt{\hat{s}_i^{(k+1)}}} g_i^{(k)}$$

0.001



# Adaptive Step Size: ADAM

(Adaptive Moment Estimation)




biased decaying momentum  
biased decaying sq. grad.


$$\begin{aligned}v^{(k+1)} &= \gamma_v v^{(k)} + (1 - \gamma_v) g^{(k)} \\s^{(k+1)} &= \gamma_s s^{(k)} + (1 - \gamma_s) (g^{(k)} \odot g^{(k)})\end{aligned}$$

# On Your Radar: ConvNets



# On Your Radar: ConvNets



# On Your Radar: ConvNets



# On Your Radar: ConvNets



# On Your Radar: Regularization

# On Your Radar: Regularization

$$\arg \min_{\boldsymbol{\theta}} \sum_{(x,y) \in \mathbf{D}} \ell(f_{\boldsymbol{\theta}}(x), y) - \beta \|\boldsymbol{\theta}\|^2$$

# On Your Radar: Regularization

$$\arg \min_{\boldsymbol{\theta}} \sum_{(x,y) \in \mathbf{D}} \ell(f_{\boldsymbol{\theta}}(x), y) - \beta \|\boldsymbol{\theta}\|^2$$

e.g. Batch norm, layer norm, dropout



# On Your Radar: Skip Connections (Resnets)

# Resources

OpenAI Spinning up