

Neural Network Function Approximation

Map of RL Algorithms

Model Based
←

Model Free
→

Learn Q	Learn π_θ
SARSA	Policy Gradient

On Policy

Q-learning

Off Policy

MLMBTRL
(Learn T, R)

Tabular

This Time

Challenges in Reinforcement Learning:


- Exploration vs Exploitation ← Bandit
- Credit Assignment ←
- Generalization ← Today

Function Approximation

Function Approximation

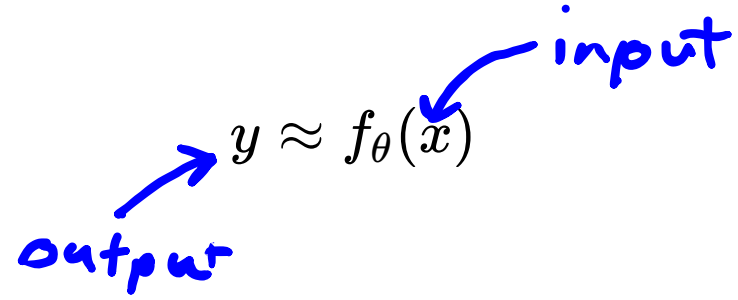
$$y \approx \underset{\tau}{f_{\theta}}(x)$$

Function Approximation

$$y \approx f_{\theta}(x)$$


input

Function Approximation



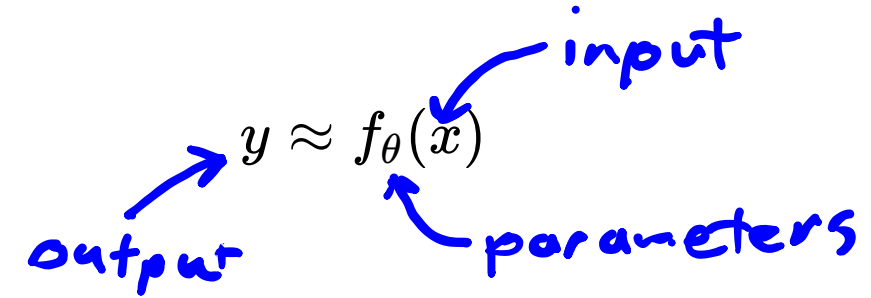
A diagram illustrating the function approximation equation $y \approx f_{\theta}(x)$. The equation is written in black. A blue arrow points from the handwritten word "output" to the variable y . Another blue arrow points from the handwritten word "input" to the variable x inside the function notation.

$$y \approx f_{\theta}(x)$$

output

input

Function Approximation



A diagram illustrating the function approximation equation $y \approx f_{\theta}(x)$. The equation is centered, with three handwritten blue arrows pointing to its components: one from the word "output" to y , one from the word "input" to x , and one from the word "parameters" to θ .

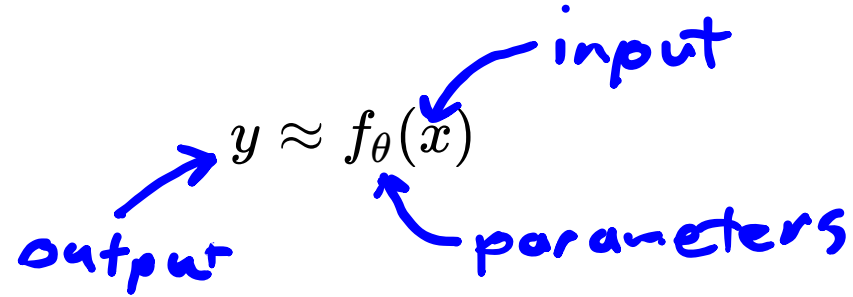
$$y \approx f_{\theta}(x)$$

output

input

parameters

Function Approximation

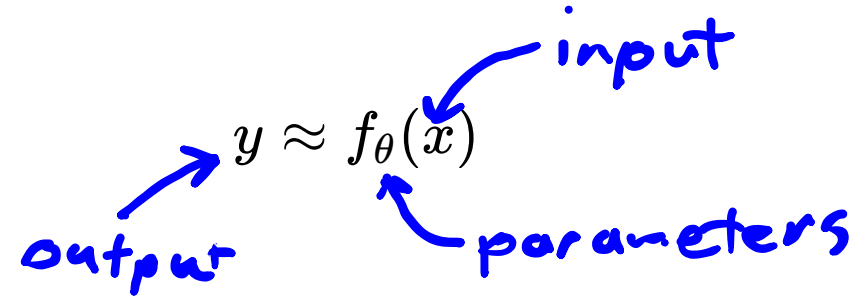


A diagram illustrating the function approximation equation $y \approx f_{\theta}(x)$. The equation is centered, with three handwritten blue annotations: an arrow pointing from the word "output" to y , an arrow pointing from the word "input" to x , and an arrow pointing from the word "parameters" to θ .

Previously, Linear:

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

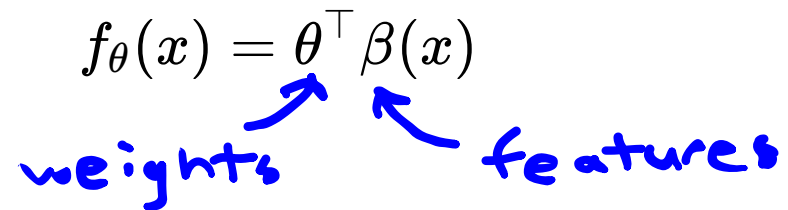
Function Approximation



A diagram showing the general function approximation equation $y \approx f_{\theta}(x)$. Three blue arrows point to the components: one from the word "output" to y , one from the word "input" to x , and one from the word "parameters" to θ .

$$y \approx f_{\theta}(x)$$

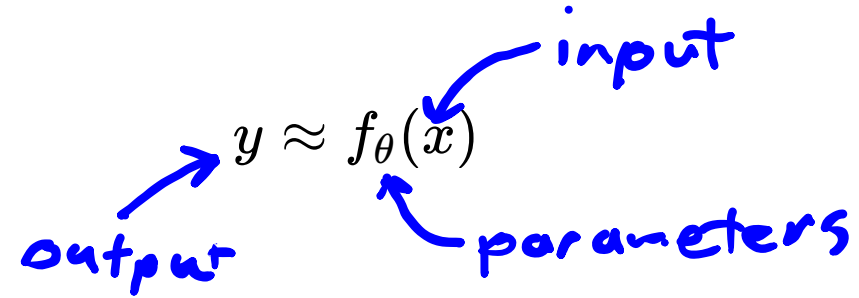
Previously, Linear:



A diagram showing the linear function approximation equation $f_{\theta}(x) = \theta^{\top} \beta(x)$. Two blue arrows point to the components: one from the word "weights" to θ , and one from the word "features" to β .

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

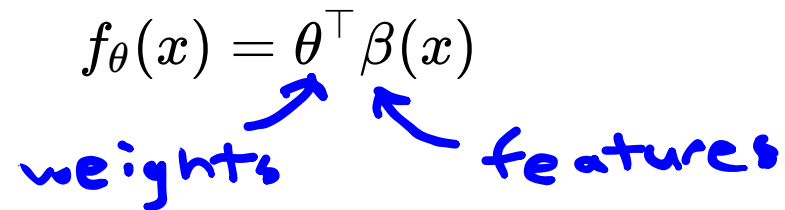
Function Approximation



A diagram showing the general function approximation equation $y \approx f_{\theta}(x)$. Three blue arrows point to the components: one from the word "output" to y , one from the word "input" to x , and one from the word "parameters" to θ .

$$y \approx f_{\theta}(x)$$

Previously, Linear:



A diagram showing the linear function approximation equation $f_{\theta}(x) = \theta^{\top} \beta(x)$. Two blue arrows point to the components: one from the word "weights" to θ , and one from the word "features" to β .

$$f_{\theta}(x) = \theta^{\top} \beta(x)$$

e.g. $\beta_i(x) = \sin(i \pi x)$

Neural Network

Neural Network



$$h(x) = \sigma(Wx + b)$$

Neural Network

$$f_{\theta}(x)$$

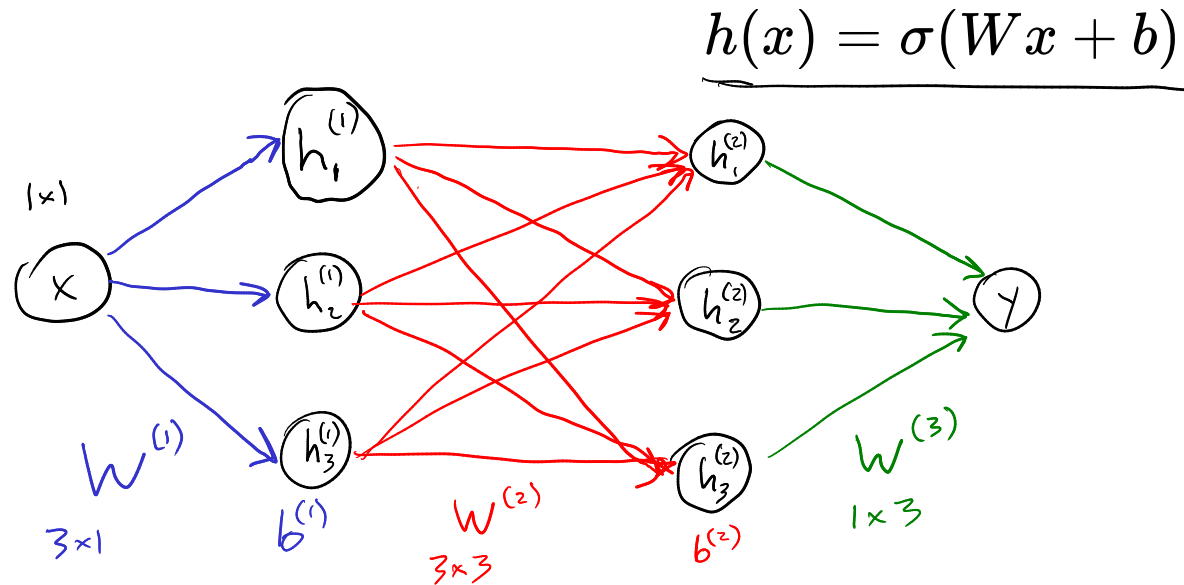
$$h_3(h_2(h_1(x)))$$

$$h(x) = \sigma(Wx + b)$$



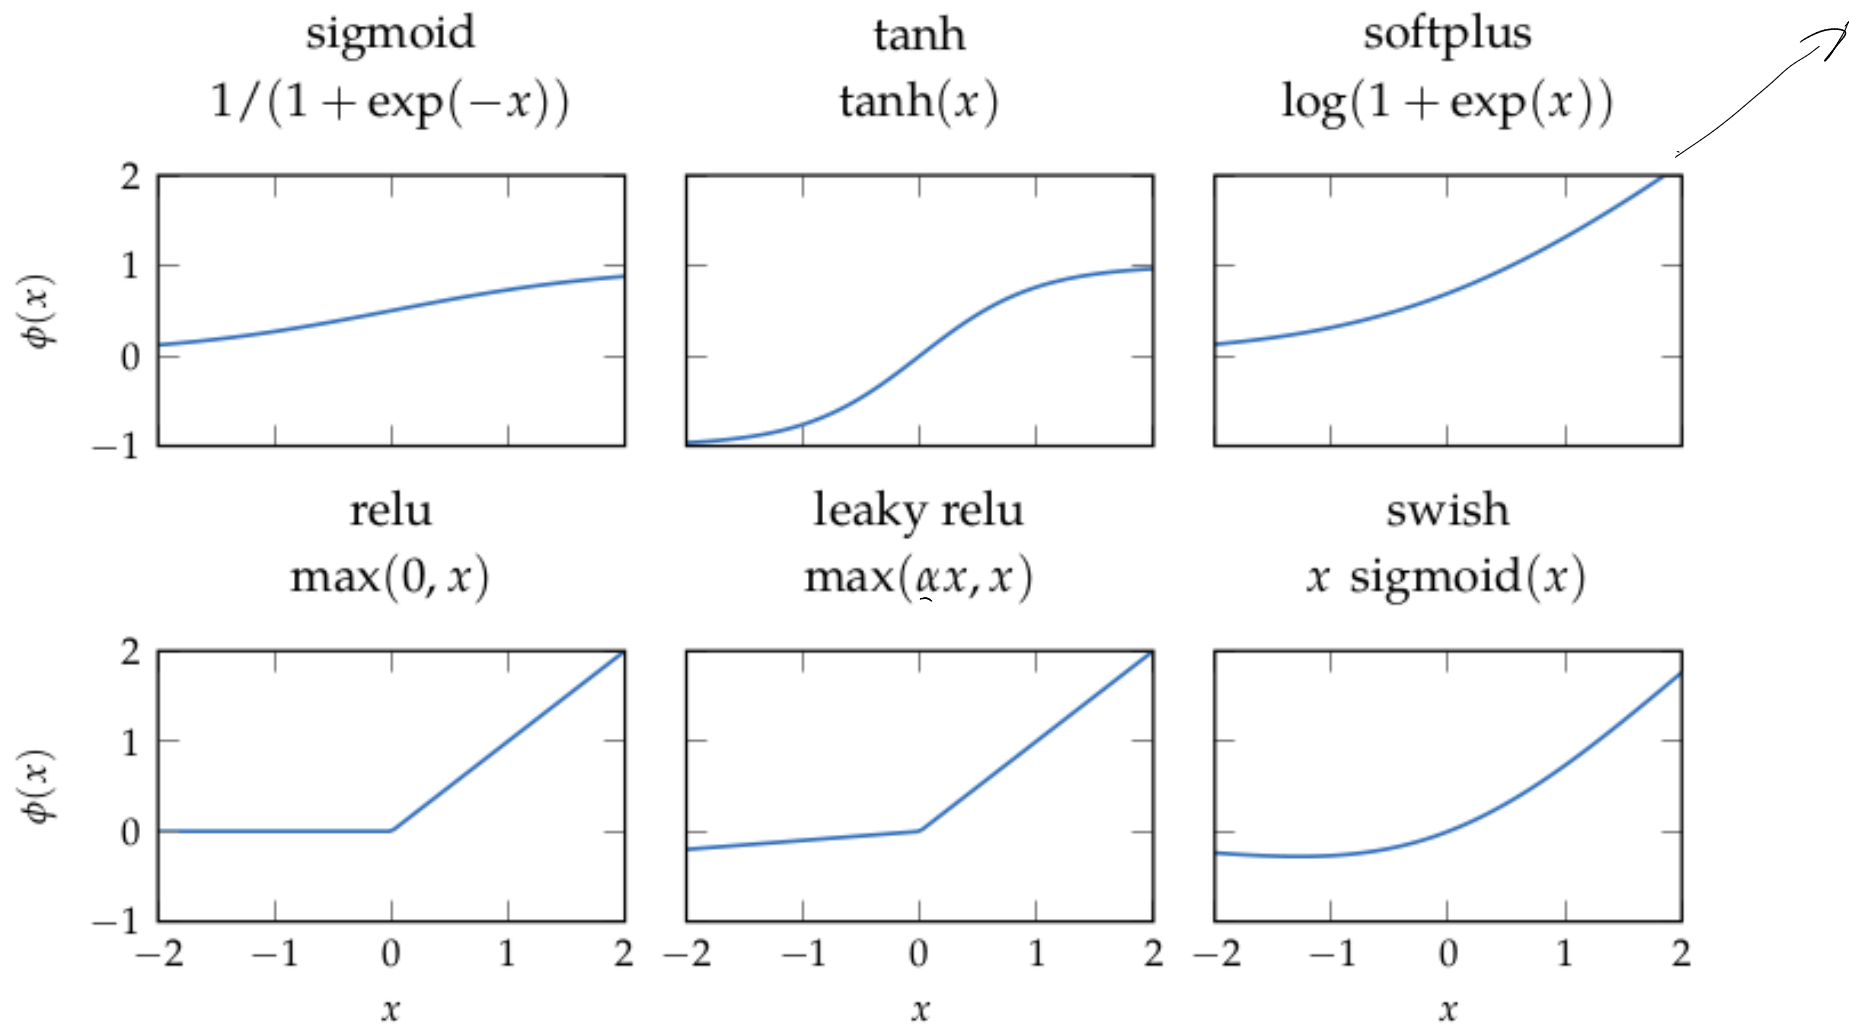
Neural Network

Neural Network




$$\underline{f_\theta(x)} = \underline{h^{(2)}(h^{(1)}(x))} = \underline{W^{(3)}} \underline{\sigma^{(2)}} \left(\underline{W^{(2)}} \underline{\sigma^{(1)}} \left(\underline{W^{(1)}} x + \underline{b^{(1)}} \right) + \underline{b^{(2)}} \right)$$

Nonlinearities




Training


Training




Training




Training



Training

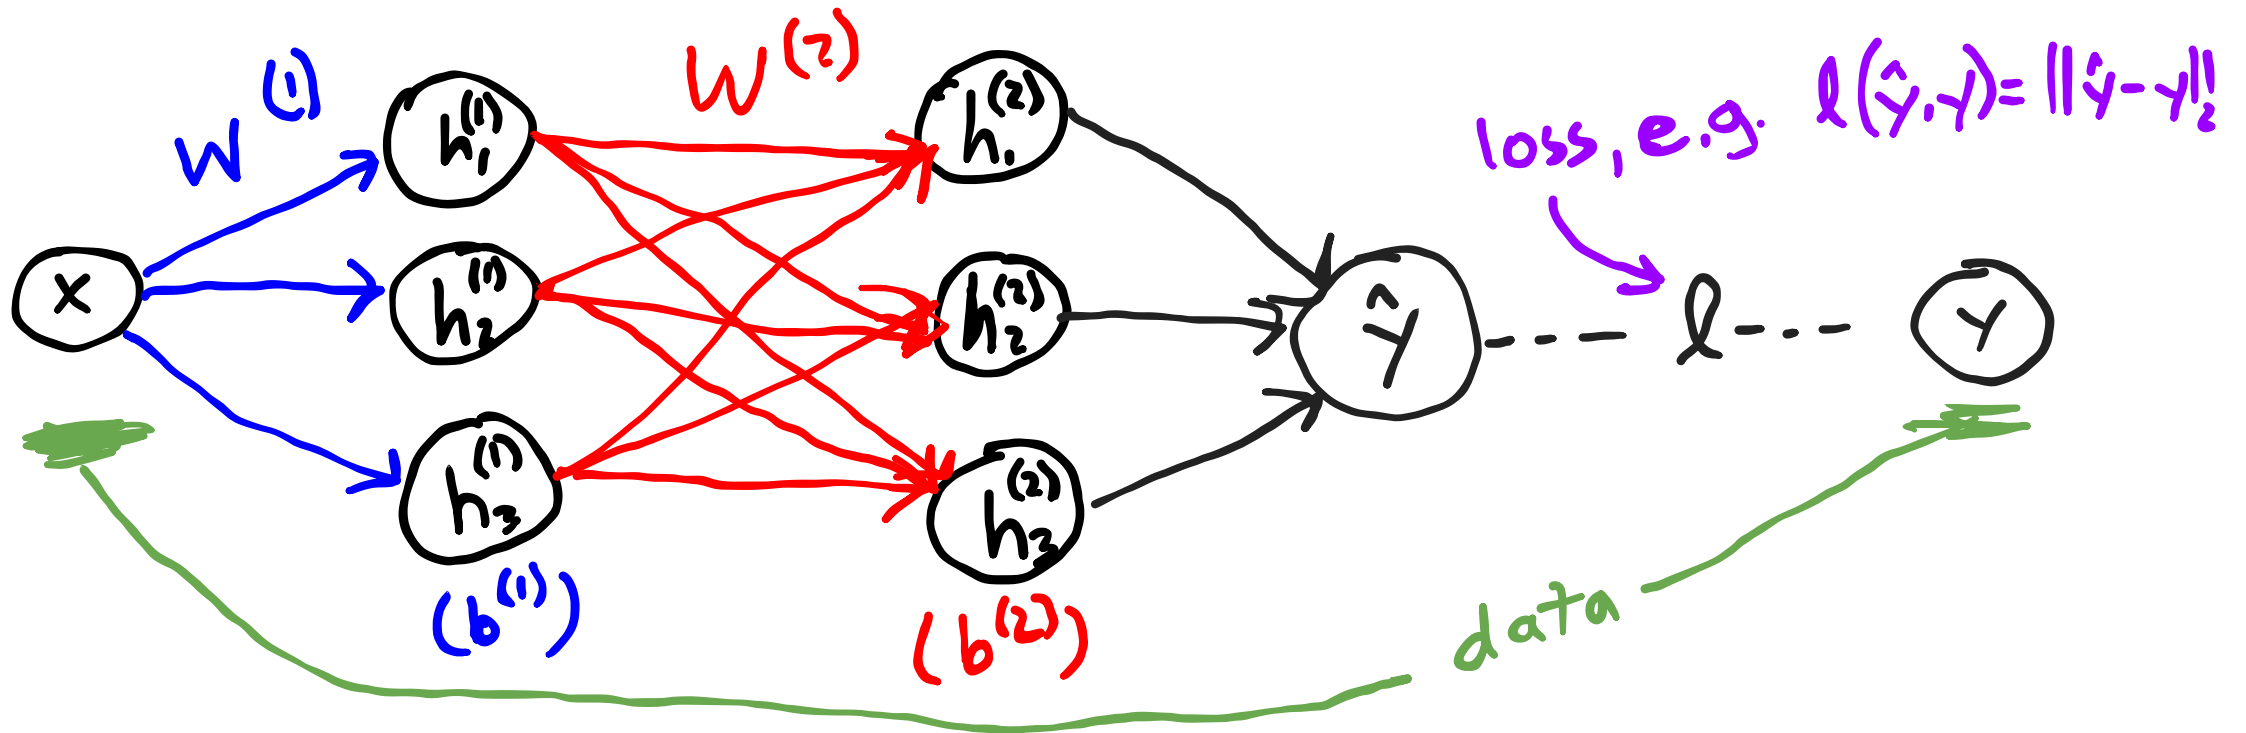


Training



$$\theta^* = \arg \min_{\theta} \sum_{(x,y) \in \mathcal{D}} l(f_{\theta}(x), y)$$

Training



$$\theta^* = \arg \min_{\theta} \sum_{(x,y) \in \mathcal{D}} l(f_{\theta}(x), y)$$

Stochastic Gradient Descent: $\theta \leftarrow \theta - \alpha \nabla_{\theta} l(f_{\theta}(x), y)$

Chain Rule

$$h_0 = h(x_0)$$

$$f \circ g \circ h = f(g(h(\cdot)))$$

$$\begin{aligned} \left. \frac{\partial f \circ g \circ h}{\partial x} \right|_{x_0} &= \left. \frac{\partial f(g(h(x)))}{\partial x} \right|_{x_0} = \left. \frac{\partial f(g(h))}{\partial h} \right|_{h_0} \left. \frac{\partial h}{\partial x} \right|_{x_0} \\ &= \left. \frac{\partial f(g)}{\partial g} \right|_{g_0} \left. \frac{\partial g(h)}{\partial h} \right|_{h_0} \left. \frac{\partial h}{\partial x} \right|_{x_0} \end{aligned}$$

$$\nabla_{\theta} l(f_{\theta}(x), y)$$

$$\begin{bmatrix} \frac{\partial l}{\partial \theta_1} \\ \vdots \\ \frac{\partial l}{\partial \theta_n} \end{bmatrix}$$

$$\theta = (w^{(1)}, b^{(1)}, w^{(2)}, b^{(2)})$$

$$l(\hat{y}, y) = \frac{1}{n} \sum_i (\hat{y}_i - y_i)^2$$

$$\hat{y} = w^{(2)} \sigma(w^{(1)} x + b^{(1)}) + b^{(2)}$$


$$\left. \frac{\partial l}{\partial w^{(2)}} \right|_{x_0} = \left. \frac{\partial l}{\partial \hat{y}} \right|_{\hat{y}_0} \left. \frac{\partial \hat{y}}{\partial w^{(2)}} \right|_{x_0} = \left. \frac{\partial l}{\partial \hat{y}} \right|_{\hat{y}_0} \left(\sigma(w^{(1)} x_0 + b^{(1)}) \right)$$

Backprop

Backprop

$$\hat{y} = \theta_2 x + \theta_1$$

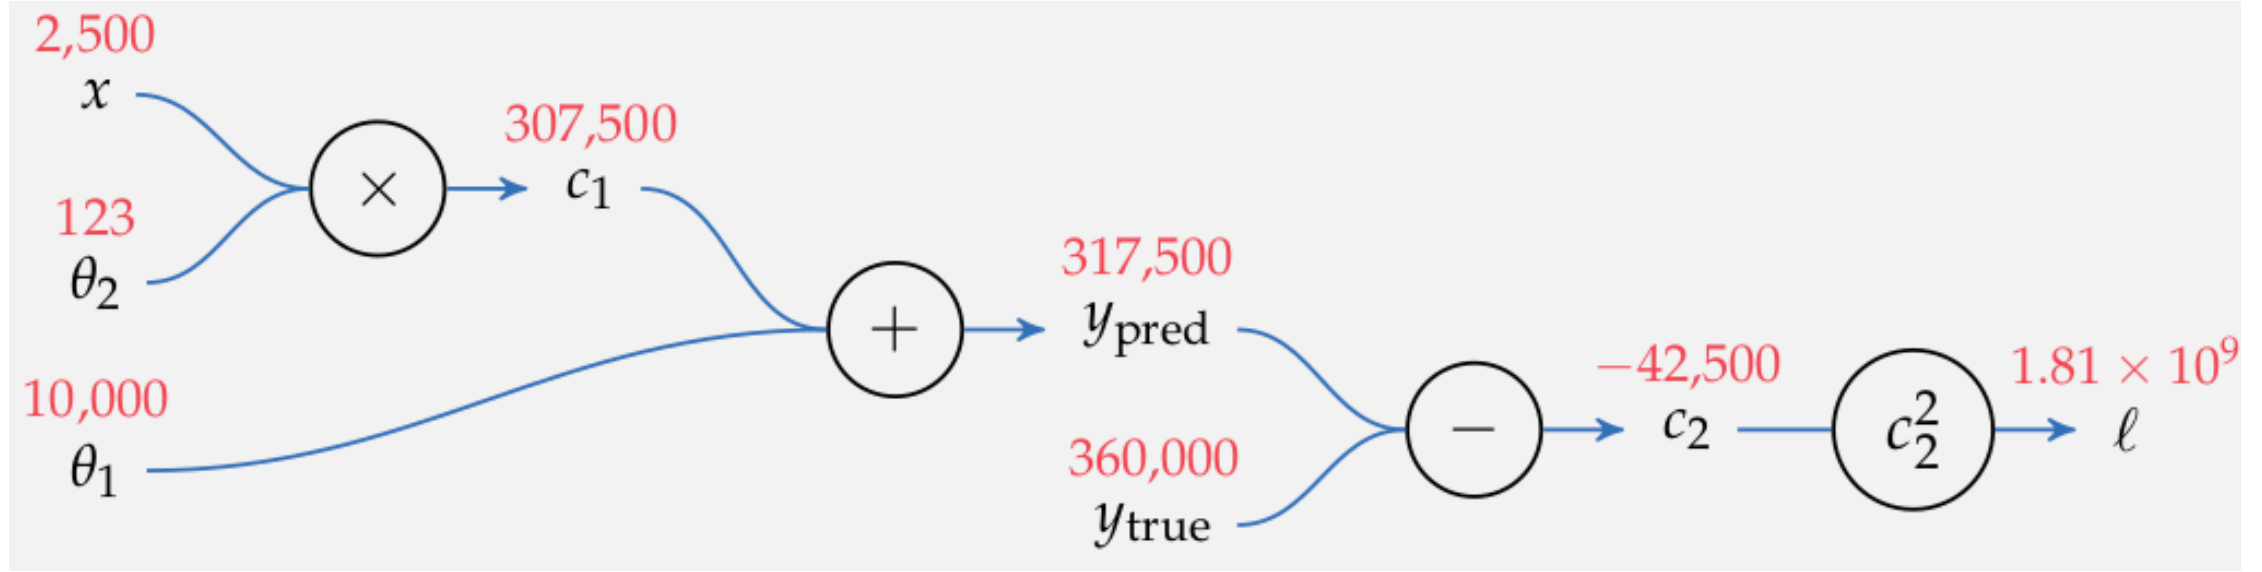
$$l(\hat{y}, y) = (\hat{y} - y)^2$$




forward: calculate values

backward: use chain rule to calculate derivatives

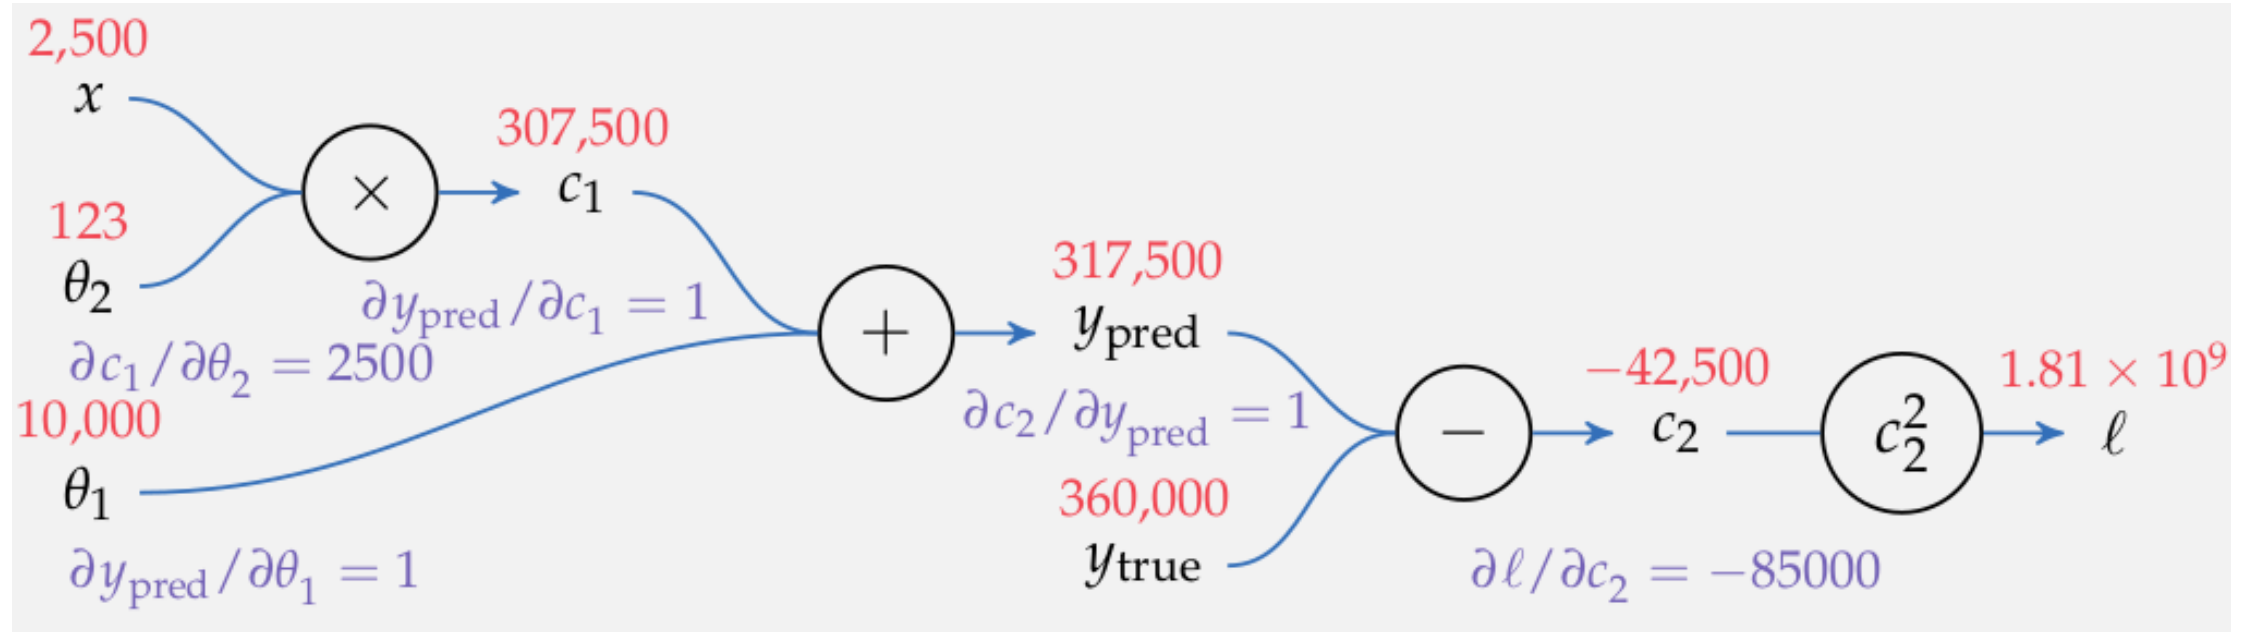
Backprop



Backprop



Backprop



backwards from ℓ to θ_i

$$\nabla_{\theta} \ell \begin{cases} \frac{\partial \ell}{\partial \theta_1} = \frac{\partial \ell}{\partial c_2} \frac{\partial c_2}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial \theta_1} = -85,000 \cdot 1 \cdot 1 = -85,000 \\ \frac{\partial \ell}{\partial \theta_2} = \frac{\partial \ell}{\partial c_2} \frac{\partial c_2}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial c_1} \frac{\partial c_1}{\partial \theta_2} = -85,000 \cdot 1 \cdot 1 \cdot 2,500 = -2.125 \times 10^8 \end{cases}$$

a “fast and furious” approach to training neural networks does not work and only leads to suffering. Now, suffering is a perfectly natural part of getting a neural network to work well, but it can be mitigated by being thorough, defensive, paranoid, and obsessed with visualizations of basically every possible thing. The qualities that in my experience correlate most strongly to success in deep learning are patience and attention to detail.

Keep calm and
lower your learning
rate

- Andrej Karpathy

Adaptive Step Size: RMSProp

SGD

$$\theta \leftarrow \theta - \alpha \underbrace{\nabla_{\theta} l(f_{\theta}(x), y)}_{g^{(k)}}$$

RMS prop

estimate
of $g \odot g$

$$\hat{s}^{(k+1)} \leftarrow \gamma \hat{s}^{(k)} + (1-\gamma) \underbrace{(g^{(k)} \odot g^{(k)})}$$

$$\theta_i^{(k+1)} = \theta_i^{(k)} - \frac{\alpha}{\epsilon + \sqrt{\hat{s}_i^{(k+1)}}} g_i^{(k)}$$

0.001

Adaptive Step Size: ADAM

(Adaptive Moment Estimation)




Figure 5.5. Gradient descent and the momentum method compared on the Rosenbrock function with $b = 100$; see appendix B.6.

biased decaying momentum $v^{(k+1)} = \gamma_v v^{(k)} + (1 - \gamma_v) g^{(k)}$

biased decaying sq. grad. $s^{(k+1)} = \gamma_s s^{(k)} + (1 - \gamma_s) (g^{(k)} \odot g^{(k)})$


corrected decaying momentum $\hat{v}^{(k+1)} = v^{(k+1)} / (1 - \gamma_v^k)$

corrected decaying sq. grad. $\hat{s}^{(k+1)} = s^{(k+1)} / (1 - \gamma_s^k)$


$$\theta^{(k+1)} \leftarrow \theta^{(k)} - \alpha \hat{v}^{(k+1)} / (\epsilon + \sqrt{\hat{s}^{(k+1)}})$$

On Your Radar: ConvNets



On Your Radar: ConvNets



On Your Radar: ConvNets



On Your Radar: ConvNets



On Your Radar: Regularization

On Your Radar: Regularization

$$\arg \min_{\boldsymbol{\theta}} \sum_{(x,y) \in \mathbf{D}} \ell(f_{\boldsymbol{\theta}}(x), y) - \beta \|\boldsymbol{\theta}\|^2$$

On Your Radar: Regularization

$$\arg \min_{\boldsymbol{\theta}} \sum_{(x,y) \in \mathbf{D}} \ell(f_{\boldsymbol{\theta}}(x), y) - \beta \|\boldsymbol{\theta}\|^2$$

e.g. Batch norm, layer norm, dropout

On Your Radar: Skip Connections (Resnets)

Resources

OpenAI Spinning up